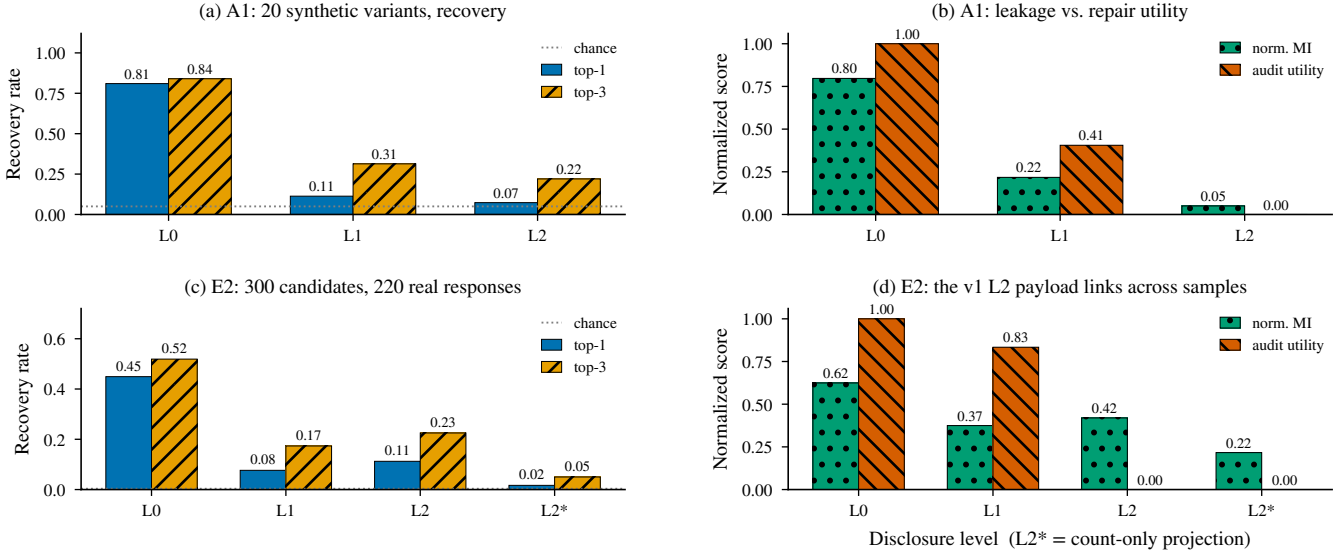

A1: policy_infer.json, 20 variants, 40 train / 15 test synthetic responses, multinomial NB; L2 is count-only. E2: policy_corpus_infer.json, 300 candidates x 220 real chat responses (19,800 test pairs), v1 certificate payload; evidence coverage 0.677. policy_hash is excluded as a feature. The v2 circuit-level L2 has NOT been re-measured; L2 and L2* are different projections of different corpora and are not averaged.